

SQL BOOTCAMP

Utku YILDIRIM, Co-Founder & Chief Product Officer at Prorize

SQL BOOTCAMP

HELLO!

AGENDA

‣ Part 1

- Introduction to Relational Databases
- Basic CREATE statements
- Basic SELECT-FROM-WHERE statements

‣ Part 2

- Complex conditions in the WHERE clause
- Data manipulation in SELECT statements

‣ Part 3

- Grouping and aggregations

‣ Part 4

- JOIN operations
- Takeaways

MOTIVATION

**Decision making process has been always demanding.
Emerging technologies take the guessing out of equation.**

**the 10-Megabyte
Computer System**

**Only
\$5995
COMPLETE**

New From IMSAI

- 10-Megabyte Hard Disk
- 5 1/4" Dual-Density Floppy Disk Back-up
- 8-Bit Microprocessor (Optional 16-bit Microprocessor)
- Memory-Mapped Video Display Board
- Disk Controller
- Standard 64K RAM (Optional 256K RAM)
- 10-Slot S-100 Motherboard
- 28-Amp Power Supply
- 12" Monitor
- Standard Intelligent 62-Key ASCII Keyboard (Optional Intelligent 86-Key ASCII Extended Keyboard)
- 132-Column Dot-Matrix Printer
- CP/M* Operating System

**You Read It Right ...
All for \$5995!**

IMSAI...Thinking ahead for the 80's

415/635-7615

Computer Division of the Fischer-Freitas Corporation
910 81st Avenue, Bldg. 14 • Oakland, CA 94621

*CP/M is a trademark of Digital Research. IMSAI is a trademark of the Fischer-Freitas Corporation

It is easier and cheaper to collect data

- Introduction of loyalty programs and online stores shortened feedback loop
- 1 GB data storage used to cost \$10 k back in 1990s vs \$ 0.1 today

Processing data is easier and faster

- Accessibility and usability of analytical tools improved drastically
- Smartphones today can process 10k more than supercomputers in 1980s

As a result:

- Customer behavior can be identified more accurately as data grows and
 - Decisions can be made to quickly adapt to an ever-changing environment
-

HOW DOES RELATIONAL DATABASES WORK?



Server-Side



Client-Side

Think of web servers

Stores data and manages connections

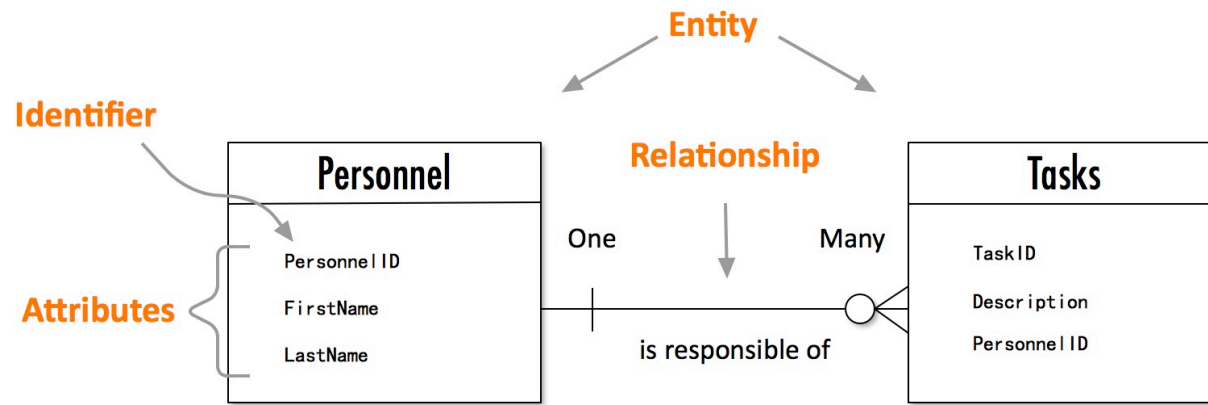
Executes instructions

Think of Chrome, Firefox, Safari, etc.

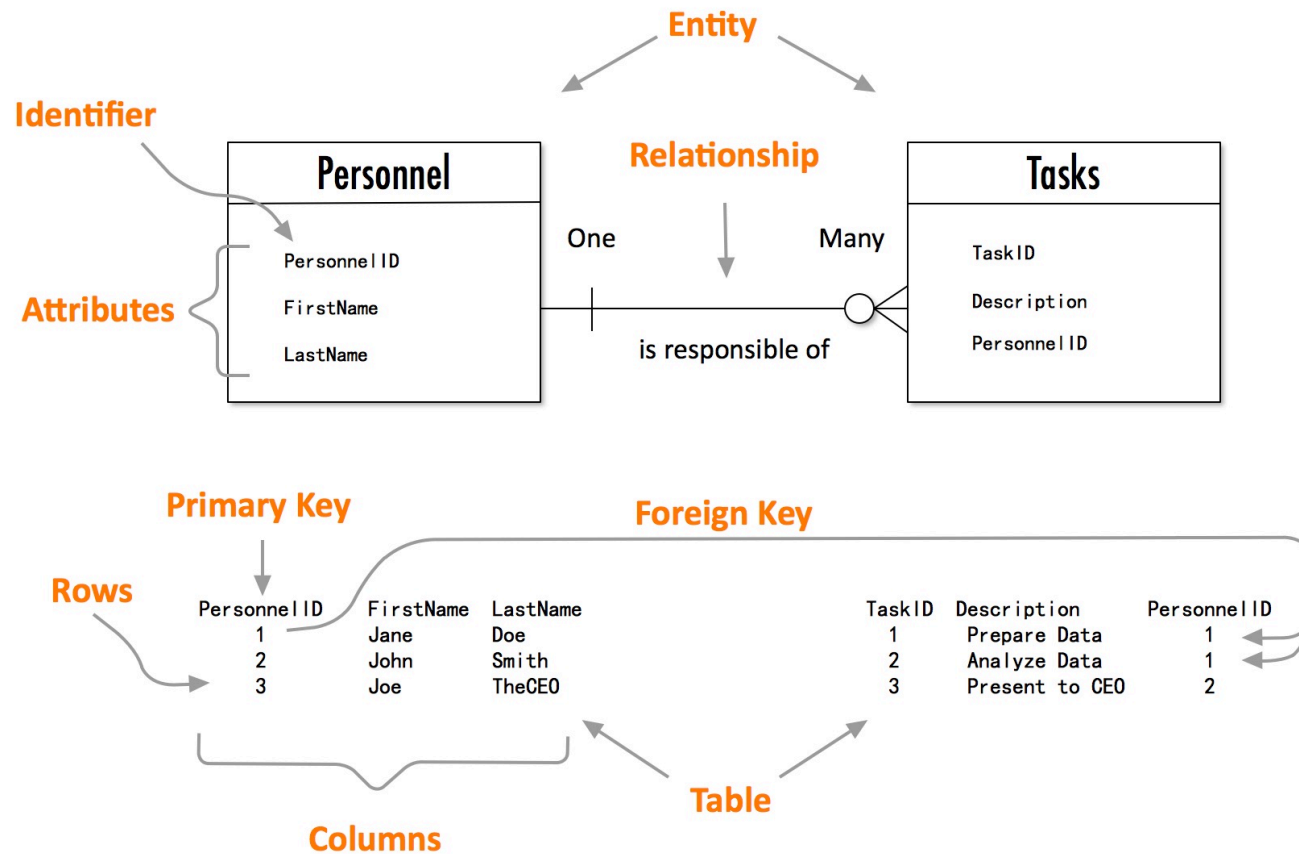
Manipulate and retrieve data

Instructions are created using Structured Query Language

DATABASE ELEMENTS





DATABASE ELEMENTS



DEDICATED CLOUD SERVER HOSTING

Rack 8



Units

Unit concept is introduced to identify each physical server and distinguish from other servers in the company. Company might track several metrics for each unit.



Configuration

A configuration represents an item in the company's product portfolio, and plays key role in the prospects decision making process.



Rentals

Prospects shop around to find the best product and price combination. If a prospect decides to rent, a new rental agreement is created




Price

Company sets and maintains a price for each configuration depending on quality, availability and competitor pricing.



WHAT IS A RELATIONAL DATABASE?



COMPONENTS OF STRUCTURED QUERY LANGUAGE

Data Definition Language (DDL)

- Business and IT teams get together and define what data needs to be collected and how it will be used
- Create, Alter, Drop, Delete, Truncate

Data Manipulation Language (DML)

- Business team uses DML to manipulate and retrieve data
- Select, Update, Insert

Data Control Language (DCL)

- IT teams decide who can access data and what they can do
 - Grant, Revoke, Roles
-

LET'S CREATE OUR TABLES

Price

```
create table price
(  
  config varchar(3),  
  startdate date,  
  enddate date,  
  rent integer,  
  constraint price_configuration_name_fk  
  foreign key (config)  
  references configuration (config)  
);
```

Configuration


```
create table configuration  
(  
  config varchar(3) primary key not null,  
  cpu integer,  
  memory integer,  
  hd integer,  
  speed integer  
);
```

Rentals

```
create table rentals  
(  
  rentid integer primary key not null,  
  unitno integer,  
  startdate date,  
  enddate date,  
  constraint rentals_unit_unitno_fk  
  foreign key (unitno)  
  references unit (unitno)  
);
```

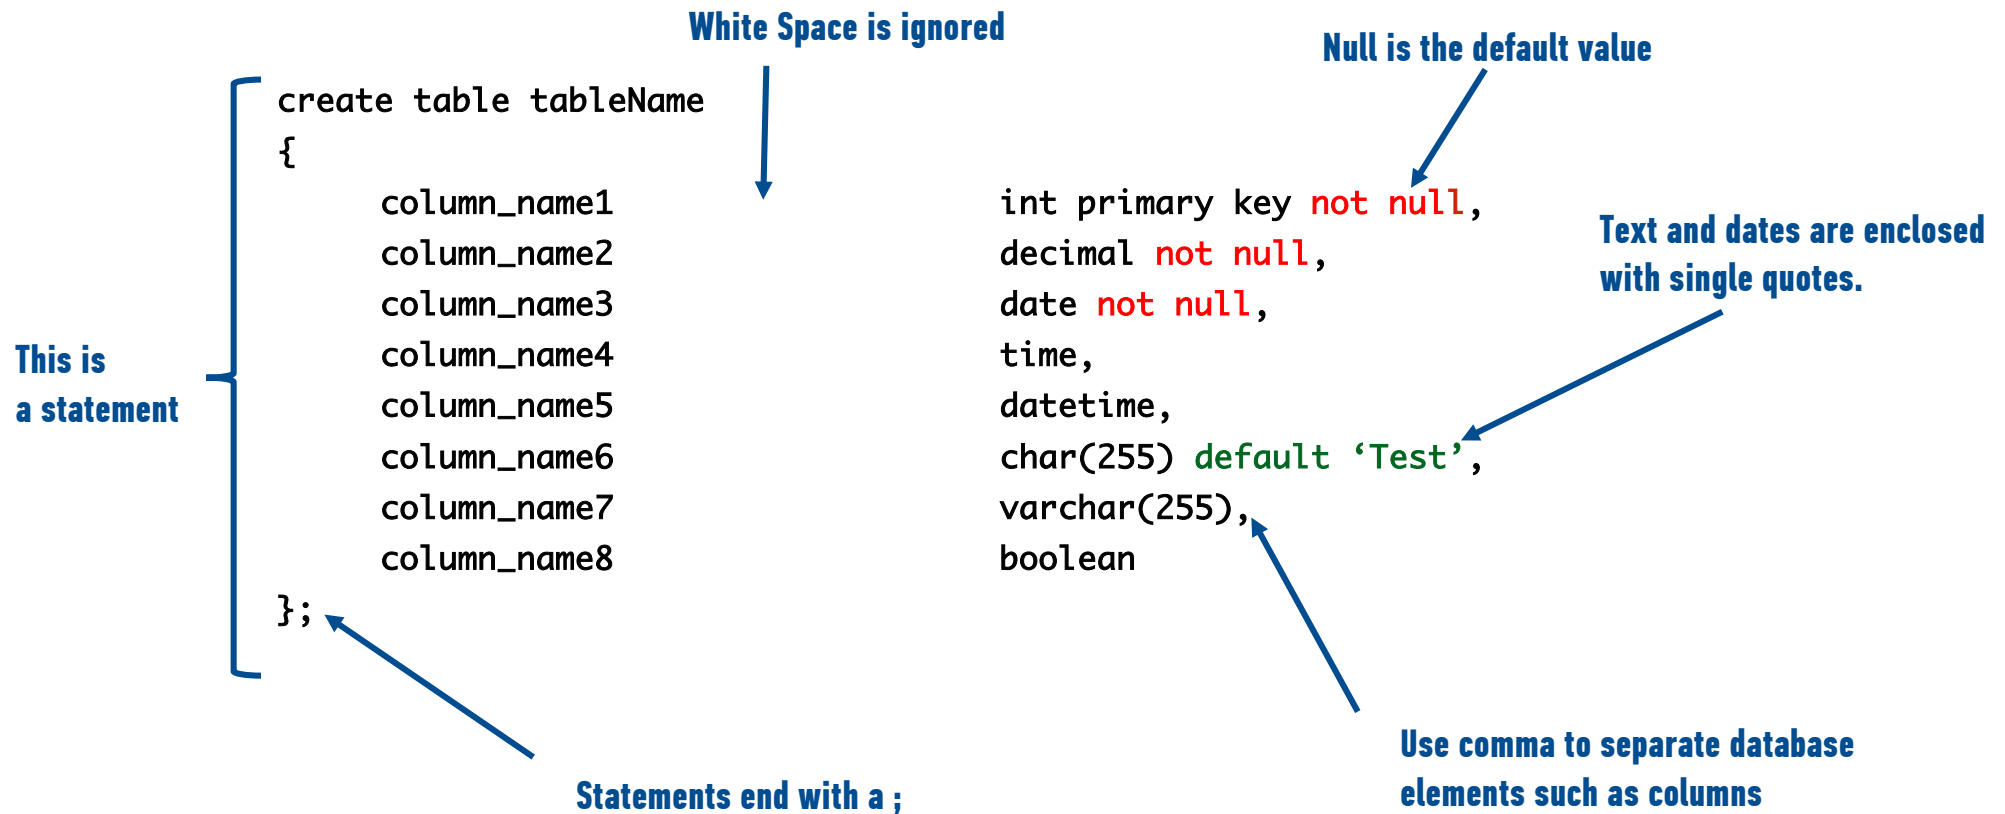
Units

```
create table unit  
(  
  unitno integer primary key not null,  
  type varchar(2),  
  config varchar(3),  
  position integer,  
  constraint unit_configuration_name_fk  
  foreign key (config)  
  references configuration (config)  
);
```



Rack 8

SYNTAX: CREATE A TABLE STATEMENT



CREATING DATABASE OBJECTS



`create database dbName`

`drop database dbName`



`create schema sName`

`drop schema sName`



`create table tableName`
`{`

`drop table tableName`



`column1 int`
`}`

SYNTAX: DATA QUERY

Select clause contains list of columns separated with comma.
All columns will be provided in the output record set

Column name is only changed in the output

A query is a statement that returns data from server as an output


```
Select column2, column3 as New  
      , column4  
      , column2 + column3 as Total  
From Table1  
Where column1 = 1;
```

White space are ignored.
It is better to use new line to increase readability

Predicates, which specify conditions that can be evaluated to SQL for each record

Expressions, which can produce either scalar values, or tables consisting of columns and rows of data. This column is calculated on the fly and only exists in the output

ALIASES



Option 1: Fully Qualified Names

```
Select DBMain.SchemaA.Table1.column1  
From DBMain.SchemaA.Table1
```

Option 2: Simplified Names and Aliases

Use DBMain

```
Select a.column2, a.column3  
From SchemaA.Table1 as a
```

```
Select b.* From Table2 as b
```

SYNTAX: FILTERING DATA

Option 1: LIMIT - OFFSET


```
Select a.column2, a.column3  
From Table1 as a  
Order by column2  
Limit {number | ALL} Offset number
```

Option 2: WHERE - DISTINCT

```
Select Distinct b.column4  
From Table2 as b  
Where b.column1 = 1
```

BRINGING ALL TOGETHER

Clauses needs to
follow this order



```
Select Distinct a.column2, a.column3  
From Table1 as a  
Where column1 = 1  
Order by column2  
Limit {number | ALL} Offset number
```

PRACTICE 1

1 - 5th sale of the January 2015

2 - Distinct configurations sold less than \$100

3 - Top three fastest configurations

4 - Cost of third most expensive unit in 'M' Type configurations on January 2016

AGENDA

‣ Part 1

- Introduction to Relational Databases
- Basic CREATE statements
- Basic SELECT-FROM-WHERE statements

‣ Part 2

- Complex conditions in the WHERE clause
- Data manipulation in SELECT statements

‣ Part 3

- Grouping and aggregations

‣ Part 4

- JOIN operations
- Takeaways

COMPARISON OPERATORS IN SQL – PART I

Where column1 = 1

=	equal to
<>, !=	is not equal to
<	less than
>	greater than
>=	greater than or equal to
<=	less than or equal to

COMPARISON OPERATORS IN SQL – PART II

IN

In a list of values listed

Color in ('Red', 'Blue', 'White')

BETWEEN

Between two dates or numbers

ClassDates between '2014-01-01' and '2016-01-01'

IS NULL

True if value of column is null

EndDate is null

LIKE

True if

'red' LIKE 'red' true

'red' LIKE 'r%' true → % indicates many characters

'red' LIKE '_e_' true → _ indicates single character

'red' LIKE 'd%' false

ARITHMETIC OPERATORS IN SQL – PART I

+	Addition
-	Subtraction
*	Multiplication
/	Division
%	Returns the integer remainder of a division.
^	Exponentiation
@	Absolute value

% , / , * have same priority and are considered before + , -
If % , / , * are used together, priority is set from left to right.

LOGICAL OPERATORS IN SQL

AND

For a row to be selected all the specified conditions must be true.

OR

For the row to be selected at least one of the conditions must be true.

NOT

For a row to be selected the specified condition must be false.

AND

Condition 1	Condition 2	SELECT
True	True	True
True	False	False
False	True	False
False	False	False

OR

Condition 1	Condition 2	SELECT
True	True	True
True	False	True
False	True	True
False	False	False

NOT is considered before AND, AND is considered before OR

EXAMPLE – COMPARISON OPERATORS

Config	CPU	Memory	HD	Speed
XS1	1	1	16	250
XS2	1	1	32	300
S1	2	1	64	325
S2	2	2	64	375
S3	2	4	128	450
M1	4	8	64	550
M2	4	8	64	650
M3	6	16	256	750
M4	6	32	512	850
L1	8	64	1024	950
L2	12	128	2048	1250

► Configurations that

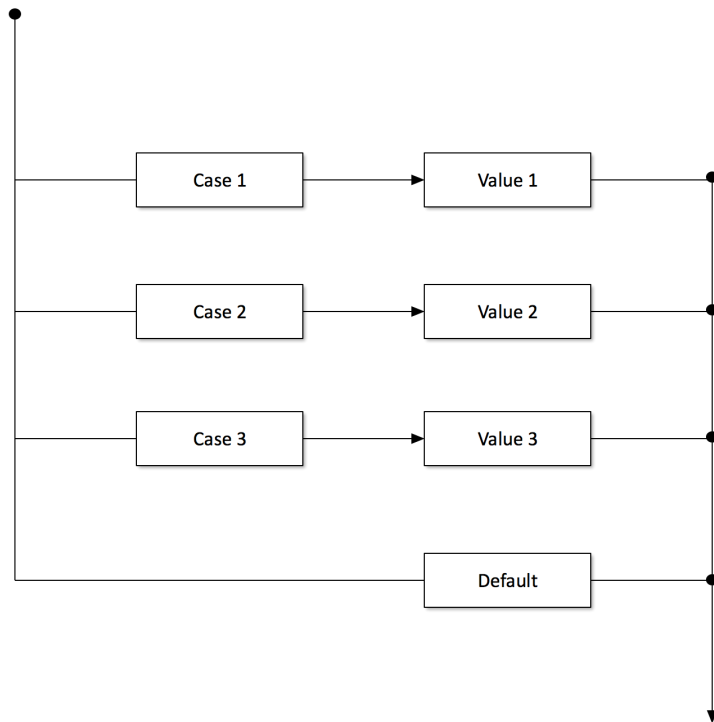
Have 'M' in its name

Have speed over 500

Memory between 32 and 128

Don't have Memory 64

CASE WHEN ... THEN ...



Select column1,

```
case when ( case 1 ) then value 1  
      when ( case 2 ) then value 2  
      when ( case 3 ) then value 3  
      else default end
```

From TableA

EXAMPLE – CASE WHEN ... THEN ...

Name	CPU	Memory	HD	Speed
XS1	1	1	16	250
XS2	1	1	32	300
S1	2	1	64	325
S2	2	2	64	375
S3	2	4	128	450
M1	4	8	64	550
M2	4	8	64	650
M3	6	16	256	750
M4	6	32	512	850
L1	8	64	1024	950
L2	12	128	2048	1250

► Create a **HDSizeGroup** column and set the value

‘EntryHD’ if HD < 65

‘MidSizeHD’ if HD > 65 but <300

‘EnterpriseHD’ if HD > 300

► Order by **HDSizeGroup**

DATE OPERATORS IN SQL

+	date '2016-04-15' + integer '7'	date '2016-04-22'
+	date '2016-04-15' + interval '7 hours'	timestamp '2016-04-15 07:00:00'
-	time '05:00' – time '03:00'	interval '02:00:00'
-	date '2016-04-22' – date '2016-04-15'	integer '7' days
date_part(text, timestamp)	date_part('hour', timestamp '2016-04-15 20:38:40')	20
date_trunc(text, timestamp)	date_part('hour', timestamp '2016-04-15 20:38:40')	timestamp '2016-04-15 20:00:00'
extract(field from timestamp)	extract(hour from timestamp '2016-04-15 20:38:40')	20
make_date(year int, month int, day int)	make_date(2016, 4, 15)	date '2016-04-15'
now()	Current date and time	
age(date startdate, date enddate)		

DATE/TIME DATA TYPES IN SQL

Date: The word followed by a date enclosed in single quotes; DATE '2006-2-18'

Other formats are allowed but this is the standard.

Time: The word followed by a time enclosed in single quotes; TIME '02:34:00'

Seconds and colons are optional.

Timestamp: Both; TIMESTAMP '2004-10-23 13:20:00'

Interval: A time interval; INTERVAL '2 days', INTERVAL '12 hours'

Many times you in PostgreSQL you don't need to type the descriptor.

For example: '2006-2-18', '02:34', etc.

PRACTICE 2

1 For each date in the rental table identify all rentals started at the beginning of a month, end of a month or in the middle of the month.

2 Identify length of total lease for each rental, and if the end date is not set then use the current date to calculate length of lease in months.

AGENDA

‣ Part 1

- Introduction to Relational Databases
- Basic CREATE statements
- Basic SELECT-FROM-WHERE statements

‣ Part 2

- Complex conditions in the WHERE clause
- Data manipulation in SELECT statements

‣ Part 3

- Grouping and aggregations

‣ Part 4

- JOIN operations
- Takeaways

AGGREGATE FUNCTIONS

▸ Aggregate functions compute summaries of data in a table

- Most aggregate functions (all except COUNT) work on a single column of numeric data
- Use an alias to name the result

▸ Aggregate functions

- COUNT: The number of rows
- SUM: The sum of the entries in a column
- AVG: The average entry in a column
- MIN, MAX: The minimum and maximum entries in a column

AGGREGATE FUNCTIONS

Name	CPU	Memory	HD	Speed
XS1	1	1	16	250
XS2	1	1	32	300
S1	2	1	64	325
S2	2	2	64	375
S3	2	4	128	450
M1	4	8	64	550
M2	4	8	64	650
M3	6	16	256	750
M4	6	32	512	850
L1	8	64	1024	950
L2	12	128	2048	1250

Select count(*) as Cnt
From Configuration

Cnt
11

Select sum(memory) as TotMem
From Configuration

TotMe
265

Select max(speed) as TopSpeed
From Configuration

TopSpeed
1250

COMBINING AGGREGATE FUNCTIONS

Name	CPU	Memory	HD	Speed
XS1	1	1	16	250
XS2	1	1	32	300
S1	2	1	64	325
S2	2	2	64	375
S3	2	4	128	450
M1	4	8	64	550
M2	4	8	64	650
M3	6	16	256	750
M4	6	32	512	850
L1	8	64	1024	950
L2	12	128	2048	1250

Select max(speed) – min(speed) as SpeedRange
From Configuration

SpeedRange
1000

PRACTICE 3

Name	CPU	Memory	HD	Speed
XS1	1	1	16	250
XS2	1	1	32	300
S1	2	1	64	325
S2	2	2	64	375
S3	2	4	128	450
M1	4	8	64	550
M2	4	8	64	650
M3	6	16	256	750
M4	6	32	512	850
L1	8	64	1024	950
L2	12	128	2048	1250

1 Find average speed per CPU count

2 Find average speed weighted by Memory

GROUP BY AGGREGATION

Name	CPU	Memory	HD	Speed
XS1	1	1	16	250
XS2	1	1	32	300
S1	2	1	64	325
S2	2	2	64	375
S3	2	4	128	450
M1	4	8	64	550
M2	4	8	64	650
M3	6	16	256	750
M4	6	32	512	850
L1	8	64	1024	950
L2	12	128	2048	1250

- › Sometimes we want to apply aggregate functions to groups of rows
- › Example: Average Memory for each CPU count.

Select columns1, aggregate functions
From tables
Where conditions
Group by columns1

GROUP BY AGGREGATION

Name	CPU	Memory	HD	Speed
XS1	1	1	16	250
XS2	1	1	32	300
S1	2	1	64	325
S2	2	2	64	375
S3	2	4	128	450
M1	4	8	64	550
M2	4	8	64	650
M3	6	16	256	750
M4	6	32	512	850
L1	8	64	1024	950
L2	12	128	2048	1250

Select CPU, avg(Memory) as avgMem
From Configuration
Group by CPU

CPU	AvgMem
1	1
2	2.33
4	8
6	24
8	64
12	128

GROUP BY AGGREGATION – FILTERING WITH HAVING

Name	CPU	Memory	HD	Speed
XS1	1	1	16	250
XS2	1	1	32	300
S1	2	1	64	325
S2	2	2	64	375
S3	2	4	128	450
M1	4	8	64	550
M2	4	8	64	650
M3	6	16	256	750
M4	6	32	512	850
L1	8	64	1024	950
L2	12	128	2048	1250

- › **HAVING** is like a **WHERE** clause, except that it applies to the results of a **GROUP BY** query
- › It can be used to select groups which satisfy a given condition

Select CPU	CPU	AvgMem
, avg(Memory) as avgMem	4	8
From Configuration	6	24
Group by CPU	8	64
Having avg(Memory) > 4	12	128

AGGREGATE FUNCTIONS

- **WHERE** refers to the rows of tables, and so cannot use aggregate functions
- **HAVING** refers to the groups of rows, and so cannot use columns which are not in the **GROUP BY**
- **Think of a query being processed as follows:**
 - Tables are combined
 - WHERE clauses
 - GROUP BY and Aggregates
 - Column selection
 - HAVING clauses
 - ORDER BY

AGENDA

‣ Part 1

- Introduction to Relational Databases
- Basic CREATE statements
- Basic SELECT-FROM-WHERE statements

‣ Part 2

- Complex conditions in the WHERE clause
- Data manipulation in SELECT statements

‣ Part 3

- Grouping and aggregations

‣ Part 4

- JOIN operations
- Takeaways

CROSS JOINS

Table 1

P1	T1
1	C1
2	C2
3	C3

Table 2

P2	T2
2	C3
3	C1
4	C2

Select *
From Table1
Cross join Table2

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

INNER JOINS – ON CONDITION WITH EQUALITY

Table 1

P1	T1
1	C1
2	C2
3	C3

Table 2

P2	T2
2	C3
3	C1
4	C2

Select *
From Table1
Inner join Table2
On Table1.P1 = Table2.P2

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

Select *
From Table1, Table2
Where Table1.P1 = Table2.P2

Select *
From Table1
Join Table2
On Table1.P1 = Table2.P2

INNER JOINS – ON CONDITION WITH EQUALITY

Table 1

P1	T1
1	C1
2	C2
3	C3

Table 2

P2	T2
2	C3
3	C1
4	C2

Select *
From Table1
Join Table2
On Table1.P1 >=
Table2.P2

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

Select *
From Table1
Join Table2
On Table1.P1 = Table2.P2

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

INNER JOINS – ON CONDITION WITH INEQUALITY

Table 1

P1	T1
1	C1
2	C2
3	C3

Table 2

P2	T2
2	C3
3	C1
4	C2

Select *
From Table1
Join Table2
On Table1.P1 < Table2.P2

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

Select *
From Table1
Join Table2
On Table1.P1 = Table2.P2

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

OUTER JOINS

Table 1

P1	T1
1	C1
2	C2
3	C3

Table 2

P2	T2
2	C3
3	C1
4	C2

Inner

Select *
From Table1
Full Outer Join Table2
On Table1.P1 = Table2.P2

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

Outer

P1	T1	P2	T2
1	C1		
2	C2		
3	C3		
		2	C3
		3	C1
		4	C2

LEFT (OUTER) JOINS

Table 1

P1	T1
1	C1
2	C2
3	C3

Table 2

P2	T2
2	C3
3	C1
4	C2

Inner

Select *
From Table1
Left Outer Join Table2
On Table1.P1 = Table2.P2

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

Left Outer

P1	T1	P2	T2
1	C1		
2	C2		
3	C3		

RIGHT (OUTER) JOINS

Table 1

P1	T1
1	C1
2	C2
3	C3

Table 2

P2	T2
2	C3
3	C1
4	C2

Inner

P1	T1	P2	T2
1	C1	2	C3
2	C2	2	C3
3	C3	2	C3
1	C1	3	C1
2	C2	3	C1
3	C3	3	C1
1	C1	4	C2
2	C2	4	C2
3	C3	4	C2

Right Outer

P1	T1	P2	T2
		2	C3
		3	C1
		4	C2

Select *
From Table1
Right Outer Join Table2
On Table1.P1 = Table2.P2



PRACTICE 5

1 Calculate total number of available units at the beginning of January 2016

2 Calculate average listed rental price that new rentals paid after September 2015

ORDER OF CLAUSES


SELECT




FROM



WHERE



GROUP BY



HAVING



ORDER BY



LIMIT

